

CUDA Path Tracing Ray Tracer

1. Repository Structure

CMake Repository (gcc/g++)

This project is a CMake-based repository that compiles with **gcc/g++** and **nvcc** (CUDA compiler). The build system supports both CPU and GPU implementations.

Organization

```
final_assignment_2400075/
|-- CMakeLists.txt          # Main CMake configuration
|-- vcpkg.json              # Dependency management
|-- src/                     # Source code directory
|   |-- main.cpp            # Application entry point
|   |-- cpu_raytracer.h     # CPU raytracer header
|   |-- cpu_raytracer.cpp   # CPU multi-threaded implementation
|   |-- gpu_raytracer.cuh   # GPU raytracer header (global memory)
|   |-- gpu_raytracer.cu    # GPU implementation (global memory)
|   |-- gpu_shared_raytracer.cuh # GPU raytracer header (shared memory)
|   |-- gpu_shared_raytracer.cu  # GPU implementation (shared memory)
|   |-- vec3.h               # 3D vector math utilities
|   |-- utility.h            # Utility functions
|   |-- interval.h           # Interval utilities
|   |-- display.cpp          # OpenGL display functions
|   |-- gui.cpp               # ImGui user interface
|   `-- helloworld.h         # Common headers and declarations
|-- build/                  # Build directory (generated)
`-- README                  # Original documentation
```

Build System The project uses **CMake 3.30+** with CUDA support. Key build features:

- **C++20** standard
- **CUDA Toolkit 12.8** integration
- **vcpkg** for dependency management (GLFW, ImGui, GLEW, GLM)
- Cross-platform support (Windows, Linux, macOS)

Build Commands:

```
# Configure
cmake -S . -B build -DCMAKE_TOOLCHAIN_FILE="$VCPKG_ROOT/scripts/buildsystems/vcpkg.cmake"

# Build
cmake --build build

# Run
./build/helloworld
```

2. Path Tracing on CUDA

Implementation Location

The CUDA path tracing implementation is located in the following files:

- a) CUDA Grid Implementation for Path Tracing File: `src/gpu_raytracer.cu` and `src/gpu_raytracer.cuh`

Key Components:

- **Kernel Launch:** `gpuPathTracingKernel` (in `gpu_raytracer.cu`)
- **Grid Configuration:** 16x16 thread blocks
- **Thread Mapping:** One thread per pixel
- **Grid Dimensions:**

```
dim3 blockSize(16, 16);
dim3 gridSize((camera.image_width + blockSize.x - 1) / blockSize.x,
               (camera.image_height + blockSize.y - 1) / blockSize.y);
```

Path Tracing Algorithm:

- **Ray Generation:** Each thread generates a camera ray for its assigned pixel
- **Ray-Sphere Intersection:** `gpuHitWorld()` and `gpuHitSphere()` functions
- **Material Evaluation:** Albedo, emission, and scattering properties
- **Hemisphere Sampling:** `gpuSampleHemisphere()` for diffuse scattering
- **Path Continuation:** Up to `maxBounces` bounces per ray
- **Sample Accumulation:** Progressive rendering with multiple samples per pixel

b) Shared Memory Details:

```
// Maximum spheres in shared memory
#define MAX_SHARED_SPHERES 64

// Cooperative loading pattern
int tid = threadIdx.y * blockDim.x + threadIdx.x;
int blockSize = blockDim.x * blockDim.y;
int spheresToLoad = (numSpheres < MAX_SHARED_SPHERES) ? numSpheres : MAX_SHARED_SPHERES;

for (int i = tid; i < spheresToLoad; i += blockSize) {
    sharedSpheres[i] = globalSpheres[i];
}
__syncthreads();
```

Kernel Launch:

```
size_t sharedMemSize = spheresToLoad * sizeof(CudaSphere);
gpuSharedPathTracingKernel<<<gridSize, blockSize, sharedMemSize>>>(...);
```

Summary:

```

CPU Raytracer initialized: 500 x 500, 25 spheres
GPU Ray Tracer (Global Memory) initialized:
    Image: 500 x 500
    Spheres: 25

GPU Ray Tracer (Shared Memory) initialized:
    Image: 500 x 500
    Spheres: 25 (max shared: 64)
    Shared memory per block: 1200 bytes

==== CUDA Path Tracer Initialized ====
Resolution: 500 x 500
CPU spheres: 25, GPU spheres: 25
Default mode: GPU (Shared Memory)

```

3. Rendered Results

CPU Implementation (Multi-threaded)

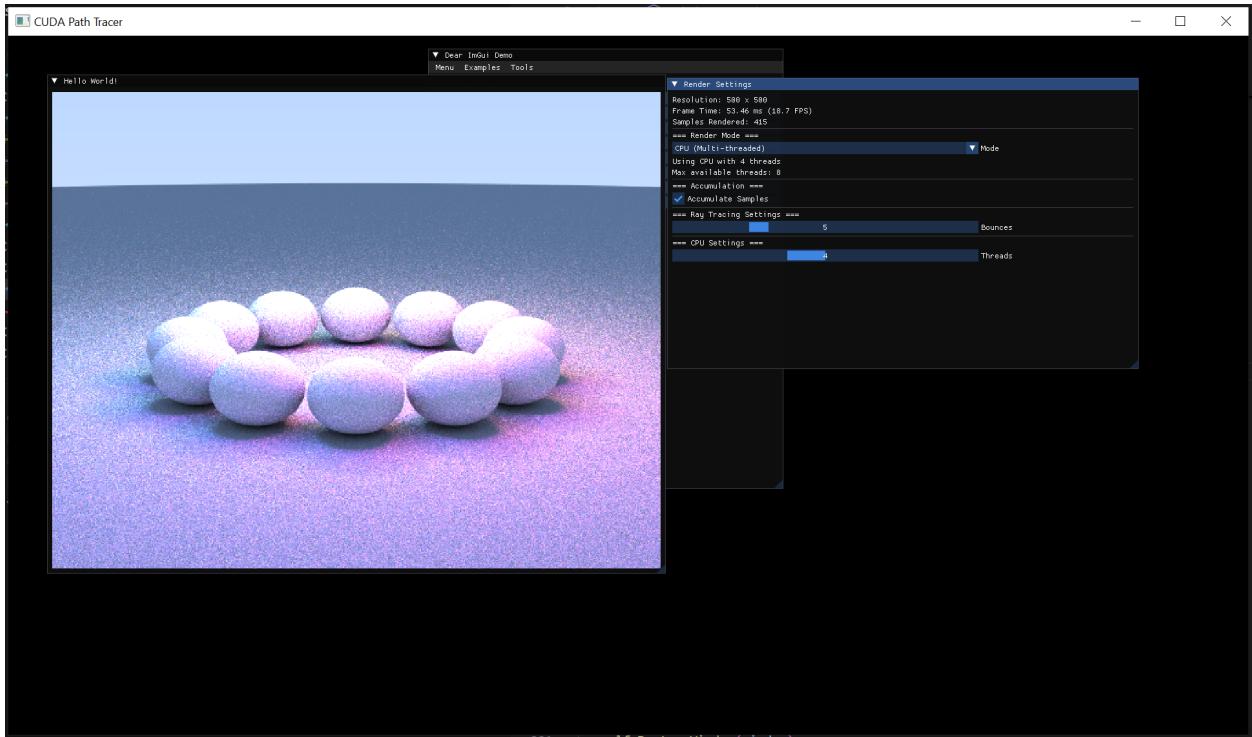


Figure 1: CPU Rendering Result

src/cpu_raytracer.cpp

GPU Global Memory Implementation

src/gpu_raytracer.cu

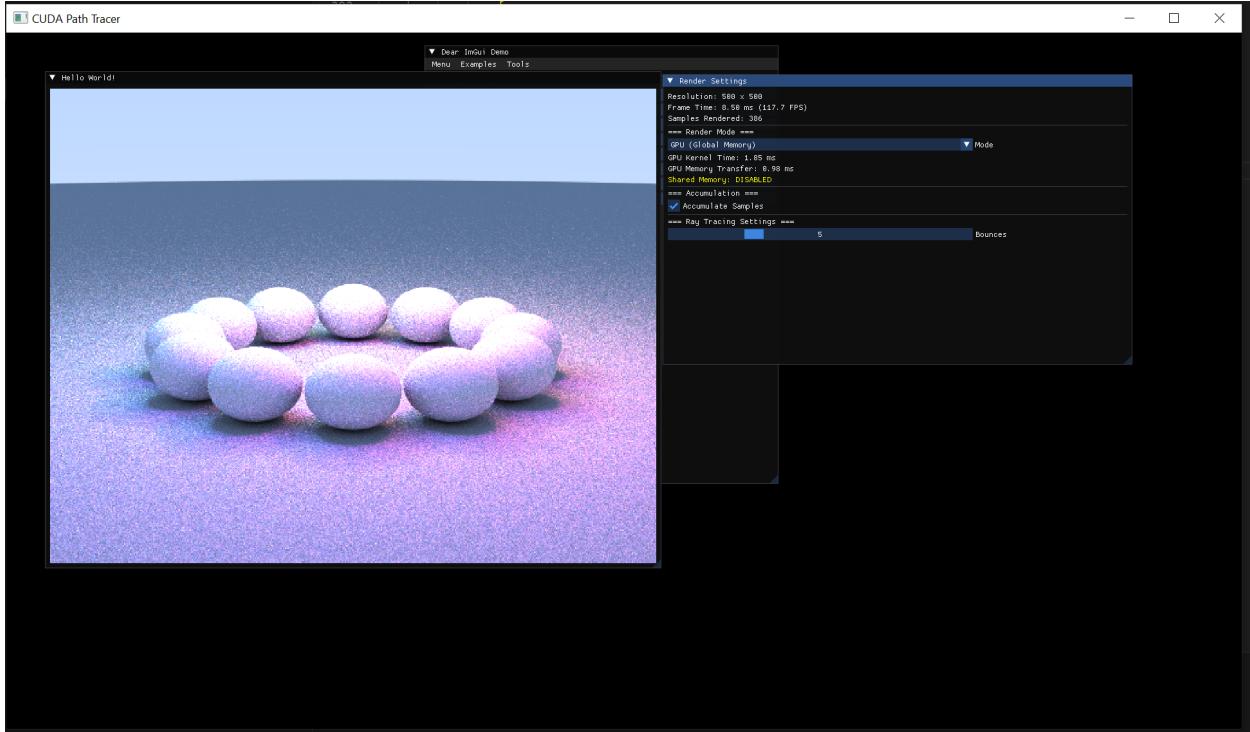


Figure 2: GPU Global Memory Rendering Result

GPU Shared Memory Implementation

`src/gpu_shared_raytracer.cu`

4. Reporting

Profiling with Nsight

Nsight Systems and Nsight Compute should be used to profile the CUDA implementations:

Profiling Commands:

```
# Profile with Nsight Systems (timeline analysis)
nsys profile --output=profile_report ./build/helloworld

# Profile with Nsight Compute (kernel analysis)
ncu --set full --export profile.ncu-rep ./build/helloworld
```

The repository includes Nsight profile reports: - CPU.nsys-rep - CPU profiling data - GPU_GLOBAL.ncu-rep - GPU global memory profiling - GPU_SHARED.ncu-rep - GPU shared memory profiling

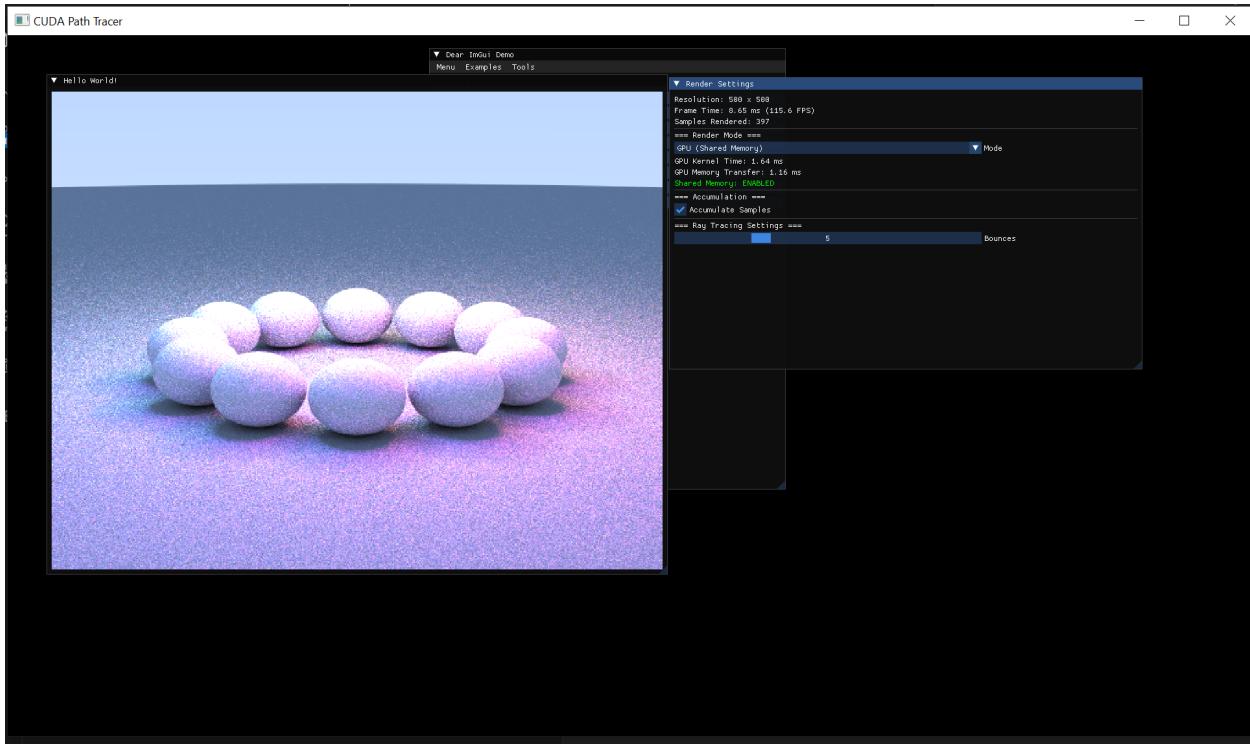


Figure 3: GPU Shared Memory Rendering Result

4. Build Instructions

Linux/macOS:

```
# Install dependencies via vcpkg
git clone https://github.com/microsoft/vcpkg.git ~/vcpkg
cd ~/vcpkg
./bootstrap-vcpkg.sh
./vcpkg integrate install

# Build project
export VCPKG_ROOT=~/vcpkg
cmake -S . -B build -DCMAKE_TOOLCHAIN_FILE="$VCPKG_ROOT/scripts/buildsystems/vcpkg.cmake"
cmake --build build

# Run
./build/helloworld
```

Windows:

```
# Install dependencies
git clone https://github.com/microsoft/vcpkg "$env:USERPROFILE\vcpkg"
cd "$env:USERPROFILE\vcpkg"
.\bootstrap-vcpkg.bat
.\vcpkg integrate install
```

```
# Build
$env:VCPKG_ROOT = "$env:USERPROFILE\vcpkg"
cmake -S . -B build -DCMAKE_TOOLCHAIN_FILE="$env:VCPKG_ROOT/scripts/buildsystems/vcpkg.cmake"
cmake --build build

# Run
.\build\Debug\helloworld.exe
```

Check GPU

```
nvidia-smi --query-gpu=name,compute_cap --format=csv
```

Set Custom CUDA Architecture

Modify CMakeLists.txt:

```
set(CMAKE_CUDA_ARCHITECTURES 61) # Change to your GPU's compute capability
```